

Assignment 3: EDA and forecast model

```
In [1]: import pandas as pd
#read the data
timeseries_data = pd.read_excel('..\data\Assignment 3 - timeseries_data.xlsx')
```

Check data types and missing values

```
In [2]: print(timeseries_data.info())
null_counts = timeseries_data.isnull().sum()
print(null_counts)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14987 entries, 0 to 14986
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   DATETIME                             14987 non-null  datetime64[ns]
1   HB_NORTH (RTLMP)                     14987 non-null  float64
2   ERCOT (WIND_RTI)                     14982 non-null  float64
3   ERCOT (GENERATION_SOLAR_RT)          14983 non-null  float64
4   ERCOT (RTLOAD)                       14987 non-null  float64
5   HOURENDING                           14987 non-null  int64
6   MARKETDAY                           14987 non-null  datetime64[ns]
7   PEAKTYPE                             14987 non-null  object
8   MONTH                               14987 non-null  object
9   YEAR                                 14987 non-null  int64
dtypes: datetime64[ns](2), float64(4), int64(2), object(2)
memory usage: 1.1+ MB
None
DATETIME                                0
HB_NORTH (RTLMP)                        0
ERCOT (WIND_RTI)                        5
ERCOT (GENERATION_SOLAR_RT)             4
ERCOT (RTLOAD)                          0
HOURENDING                             0
MARKETDAY                              0
PEAKTYPE                               0
MONTH                                  0
YEAR                                   0
dtype: int64
```

```
In [3]: means = timeseries_data.mean(numeric_only=True)

# Fill missing values with the mean of each column
timeseries_data.fillna(means, inplace=True)
```

Descriptive statistics

```
In [4]: print(timeseries_data.describe())
```

	HB_NORTH (RTLMP)	ERCOT (WIND_RTI)	ERCOT (GENERATION_SOLAR_RT) \
count	14987.000000	14987.000000	14987.000000
mean	25.766417	7532.436283	291.989714
std	46.361945	3992.218676	370.865092
min	-17.860000	54.440000	0.000000
25%	18.041250	4138.390000	0.000000
50%	20.057500	7283.460000	22.200000
75%	25.030000	10851.280000	608.580000
max	2809.357500	20350.400000	1257.540000

	ERCOT (RTLOAD)	HOURENDING	YEAR
count	14987.000000	14987.000000	14987.000000
mean	42371.673703	12.495763	2017.415493
std	9874.339631	6.922309	0.492823
min	25566.511248	1.000000	2017.000000
25%	35431.636526	6.000000	2017.000000
50%	39934.007113	12.000000	2017.000000
75%	47873.100786	18.000000	2018.000000
max	73264.662123	24.000000	2018.000000


RTLMP has an extra maximum value, which could indicate an unusually high outlier or a data error. Locational marginal prices can sometimes be negative due to an oversupply of electricity or minimal demand. However, an unexpected maximum value might also suggest data errors or specific market conditions

Similarly, WIND_RTI has an extra minimum value. Like the previous example, this unusually low value could be an outlier or an error.

Visualization of RTLMP over time

```
In [5]: import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt

# Visualization of RTLMP over time
plt.figure(figsize=(14, 7))
plt.plot(timeseries_data['DATE TIME'], timeseries_data['HB_NORTH (RTLMP)'], label='RTLMP')
plt.title('RTLMP over Time')
plt.xlabel('Date')
plt.ylabel('RTLMP ($/MWh)')
plt.legend()
plt.show()
```

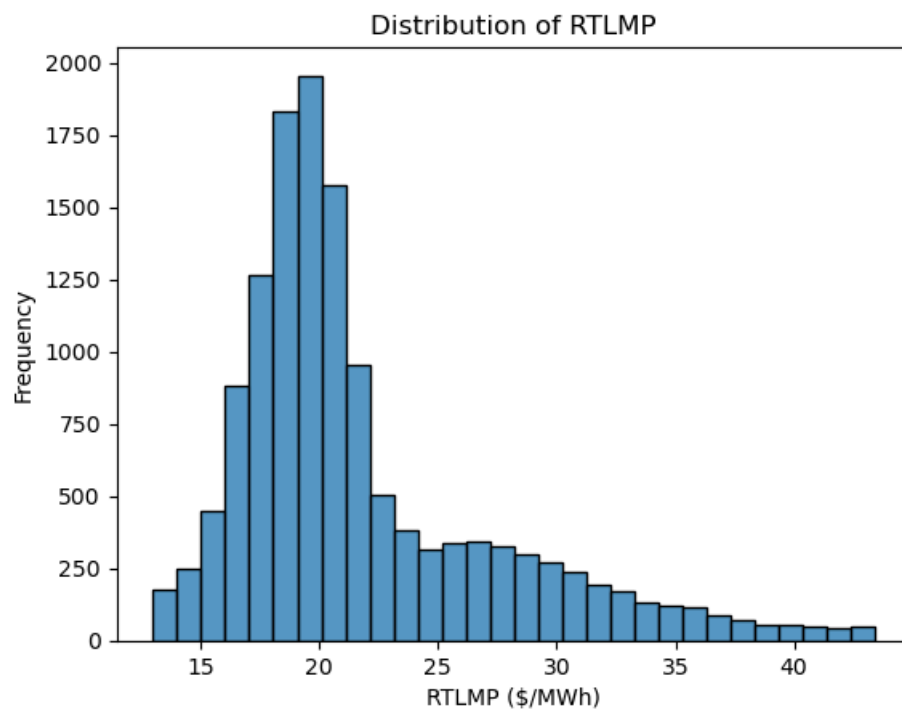


This graph depicts RTLMP over time. Most of the data points are below 500 USD/MWh. There are several noticeable spikes where the RTLMP reaches particularly high values. The most significant spike occurs around February 2018, where it exceeds 2500 USD/MWh. Other spike occurs around June 2018. These peaks suggest periods of high price volatility, which could be influenced by various factors such as demand decreases, supply shortages, or specific market events

```
In [6]: # Distribution of RTLMP

range_l = timeseries_data['HB_NORTH (RTLMP)'].quantile(0.05)
range_r = timeseries_data['HB_NORTH (RTLMP)'].quantile(0.95)

# Plot the histogram with range between 5% and 95%
sns.histplot(timeseries_data['HB_NORTH (RTLMP)'], bins=30,binrange=(range_l, range_r))
plt.title('Distribution of RTLMP')
plt.xlabel('RTLMP ($/MWh)')
plt.ylabel('Frequency')
plt.show()
```



The histogram displays the distribution of RTLMP values. The distribution shows a right skew. It means that there are a few instances of higher prices while most RTLMP values around the common price range.

Feature Engineering

```
In [7]: timeseries_data['weekday'] = timeseries_data['DATETIME'].dt.weekday
timeseries_data['hour'] = timeseries_data['DATETIME'].dt.hour
timeseries_data['month'] = timeseries_data['DATETIME'].dt.month
timeseries_data['day'] = timeseries_data['DATETIME'].dt.day
timeseries_data['PEAKTYPE'] = timeseries_data['PEAKTYPE'].map({'OFFPEAK': 0, 'WEPEAK': 1, 'WDPEAK': 2})
timeseries_data.head()
```

*#Adding features like weekday, hour, and month helps capture daily, weekly, and seasonal patterns in RTLMP.
#PeakType is 1 if weekday peak, 0 if offpeak, 2 if weekend peak*


```
Out[7]:
```

	DATETIME	HB_NORTH (RTLMP)	ERCOT (WIND_RTI)	ERCOT (GENERATION_SOLAR_RT)	ERCOT (RTLLOAD)	HOURENDING	MARKETDAY	PEAKTYPE	MONTH	YEAR	wee
0	2017-01-01 01:00:00	23.3575	2155.31	0.0	29485.791355	1	2017-01-01	0	JANUARY	2017	
1	2017-01-01 02:00:00	21.4650	2313.81	0.0	28911.565913	2	2017-01-01	0	JANUARY	2017	
2	2017-01-01 03:00:00	20.7350	2587.68	0.0	28238.258175	3	2017-01-01	0	JANUARY	2017	
3	2017-01-01 04:00:00	20.2700	2748.65	0.0	27821.000513	4	2017-01-01	0	JANUARY	2017	
4	2017-01-01 05:00:00	20.1200	2757.49	0.0	27646.942413	5	2017-01-01	0	JANUARY	2017	

```
In [8]: # Correlation matrix heatmap
plt.figure(figsize=(6, 4))
sns.heatmap(timeseries_data.corr(), annot=True, fmt=".2f", cmap='coolwarm')
plt.show()
```

C:\Users\02221\AppData\Local\Temp\ipykernel_23936\1983811340.py:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(timeseries_data.corr(), annot=True, fmt=".2f", cmap='coolwarm')
```



This correlation matrix provides insights into the relationships between the dependent variable (RTLMP) and other variables.

There is a moderate negative correlation (-0.15) between RTLMP and WIND_RTI. This suggests that higher wind energy production might lead to lower locational marginal prices, possibly due to increased supply lowering energy costs. RTLMP shows a positive correlation (0.15) with GENERATION_SOLAR_RT. This could indicate that increased solar power generation corresponds to higher marginal prices. A positive correlation (0.24) exists between RTLMP and RTLOAD, indicating that higher energy demand is associated with higher locational marginal prices; We can find a positive correlation (0.11) between RTLMP and the peak type. On weekday or weekend peak hour, the RTLMP tends to be higher; There is a positive correlation (0.11) between RTLMP and the hour of the day. This implies that RTLMP tends to be higher at certain times of the day. For example energy usage increases during peak demand that leads to higher prices.

Prediction

```
In [9]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import OneHotEncoder

X = timeseries_data.drop(['HB_NORTH (RTLMP)', 'PEAKTYPE', 'MONTH', 'DATETIME', 'MARKETDAY'], axis=1) #independent var
y = timeseries_data['HB_NORTH (RTLMP)'] # dependent variable
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Train model
model = RandomForestRegressor()
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)
print(f'RMSE: {mean_squared_error(y_test, y_pred, squared=False)}') # RMSE

RMSE: 27.807612057627164
```

```
In [10]: timeseries_data['pred'] = model.predict(X)

# Visualization of RTLMP over time
plt.figure(figsize=(14, 7))
plt.plot(timeseries_data['DATETIME'], timeseries_data['HB_NORTH (RTLMP)'], label='RTLMP')
plt.plot(timeseries_data['DATETIME'], timeseries_data['pred'], label='pred')
plt.title('RTLMP over Time')
plt.xlabel('Date')
plt.ylabel('RTLMP ($/MWh)')
plt.legend()
plt.show()
```



Note:for assignment3, I made an independent python file named forecast.py under the "code" folder

In []:

In []: